

Andrew Ma

Lab 5

CPE 435

2/17/21

## Theory

This lab is for exploring message queues. Message queues are one form of IPC (inter-process communication), and messages can be passed between processes using message queues. All the messages are stored in the kernel, and each message has an associated message queue id (msqid). The msg queue structure (msqid\_ds) is in <sys/msg.h>, and the user defines the msg structure. Each msg has a long messageType and a data portion, which can be any size. The msqid\_ds structure has the PID of last sent and received messages, the max and current number of bytes allowed on the queue, the time of last sent, received, and changed messages, and a struct ipc\_perm describing the operation permissions.

The int msgget(key\_t key, int msgflg) function takes in a key to associate with a message queue, and it can either create a new message queue or get an existing message queue. The return value is the newly created or existing message queue identifier (msqid) that can be used in msgsnd, msgrcv, and msgctl. If the msgflg is set to IPC\_CREAT and no message queue with the same key exists, then a new message queue is created. The permissions for the message queue can be set by ORing the permissions number with IPC\_CREAT.

The int msgsnd(int msqid, void const\* msgp, size\_t msgsz, int msgflg) function is for sending messages onto the message queue specified by msqid. The msgp pointer points to a user defined message structure that has a long mtype (specifies type of message and must be a positive number), and then a data portion that holds the data bytes of the message. The msgsz is the size of the data portion not the size of the struct.

The msgflg sets the action to take if the number of bytes on the queue has reached the maximum number of bytes on the queue, or we have reached the system's maximum number of messages. If the IPC\_NOWAIT is set, the message is not sent and the calling thread returns immediately. If the IPC\_NOWAIT is not set, then the calling thread suspends execution until 1) the flag is no longer 0 which will then send the message, 2) the message queue identifier is removed which will cause an error and return -1, or 3) the calling thread receives a signal causing the message to not be sent and resuming execution in sigaction.

If the message is successfully sent, the message queue number adds 1, the message queue last PID is set to the caller's PID, and the message queue last set time is set to the current time, and it returns 0. If it fails, it returns -1.

The `msgrcv(int msqid, void* msgp, size_t msgsz, long msgtyp, int msgflg)` function is for receiving messages. It reads a message from the specified `msqid` queue, and places the message into a buffer `msgp` that the user has created, and this buffer can be cast to the message struct type to access the member variables by name. The `msgsz` is the size in bytes of the message data (not the size of the message struct). The `long msgtyp` is the type of message, and if the type number is 0, we get the first message on the queue. If the type number is greater than 0, then we get the first message on the queue with that type number.

The `msgflg` sets the action to take if the message type is not on the queue. If `IPC_NOWAIT` is set, the calling thread returns -1 immediately and sets `errno=ENOMSG`. If `IPC_NOWAIT` is not set, then the calling thread will be suspended until 1) a message of the desired type is placed on the queue (waits for message of that type), 2) the message queue is removed and `errno=EIDRM` and -1 is returned, or 3) the calling thread receives a signal and the message will not be received and the calling thread will resume according to `sigaction`.

If the message is successfully received, the message queue number subtracts 1, the last read PID is set to the caller's PID, and the last read time is set to the current time. If it is successful, it will return the number of bytes actually placed into the buffer *mtext*. If it fails, it will return -1;

The `msgctl(int msqid, int cmd, struct msqid_ds* buf)` function is for performing a control operation on a message queue specified by `msqid`. If the `struct msqid_ds* buf` is not NULL, then it points to a `msqid_ds` structure with the properties of the message queue. Some `cmd` operations are `IPC_STATE` (read the message queue's `msqid_ds` values into the local pointer `buf`), `IPC_SET` (updates the message queue's member variables with the ones in the local pointer `buf`), and `IPC_RMID` (immediately removes the message queue and awakens all waiting reader and writer processes with an error return and `errno=EIDRM`).

## Observations

## Process A sends message to process B and waits

```

u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ gcc processB.c -o processB
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ ./processA
A > Hi I am processA
Waiting ...
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ ./processB
Waiting ...
A: "Hi I am processA"
B >
I

```

## Process B sends message to process A and waits

```

u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ gcc processB.c -o processB
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ ./processA
A > Hi I am processA
Waiting ...
B: "I am pB"
A >
I
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ ./processB
Waiting ...
A: "Hi I am processA"
B > I am pB
Waiting ...

```

## Process A sends “Exit”, causing both to exit

```

u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ gcc processB.c -o processB
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ ./processA
A > Hi I am processA
Waiting ...
B: "I am pB"
A > I like
Waiting ...
B: "pears"
A > Exit
Closed Message Queue Successfully
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ ./processB
Waiting ...
A: "Hi I am processA"
B > I am pB
Waiting ...
A: "I like"
B > pears
Waiting ...
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$

```

## Process B sends “Exit”, causing both to exit

```

u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ ./processA
A > i am
Waiting ...
B: 'bob'
A > test
Waiting ...
Closed Message Queue Successfully
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ ./processB
Waiting ...
A: "i am"
B > bob
Waiting ...
A: "test"
B > Exit
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$

```

If process B runs before process A creates a message queue

```

u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$ ./processB
Failed to get existing message queue
u@-:/mnt/c/Users/U/Documents/SCHOOL/OS_Labs/Lab_05$

```

## Lab Questions

- a. How do you make a process wait to receive a message and not return immediately?

In the `msgrcv()` function, 0 can be passed in as the flag to wait to receive a message. To return immediately, `IPC_NOWAIT` is passed in as the flag.

- b. Message Queue vs Shared Memory (use and differences)

Message Queues and Shared Memory are both used for IPC to share information between processes. Shared Memory can be accessed with random access, so it is very fast. Message Queues can be accessed with linear access since it is a linked list of messages. Both of them must be created with a `{shm,msg}get()` function and both can be destroyed with `{shm,msg}ctl()` with the `IPC_RMID` flag. Shared Memory can use a local pointer that points to the shared memory address, and use the local pointer normally to read and write data. But, it must detach the local pointer to the shared memory address when the process ends. Message Queues must use `msgsnd()` and `msgrcv()` with a msg buffer to send and receive data. There is no attached pointer in message queues, so there is no detach function for message queues. Another main difference is the synchronization. Shared memory has no synchronization provided, so there can be a flag in the shared structure that different processes can use to take or receive control; but, all the logic to handle the flag variable will have to be defined by each process. In Message Queues, synchronization is provided, and the `IPC_NOWAIT` flag can be set for both sending or receiving to not wait, or the flag can be 0 to wait until a message is able to be sent or received.

- c. Use of function `ftok()`, what is its use?

The `ftok()` function converts a pathname and a project identifier to an IPC key. It takes in a filename, and a project id integer to generate a `key_t` type IPC key that can be used with `msgget()`, `semget()`, or `shmget()`. It will generate the same key if the same filename and same project id integer is

used. In my project I set the key\_t KEY=12345, but ftok() can be used to generate a key automatically.

d. What does IPC\_NOWAIT do?

In the msgsnd() function, the IPC\_NOWAIT flag will return immediately if the number of bytes on the queue has reached the max number of bytes on the queue, or if we have reached the maximum number of messages. In the msgrcv() function, the IPC\_NOWAIT flag will return immediately if the requested message type (if message type is 0, then it means all message types) is not on the queue.

## Conclusion

The program worked as expected, and the processes could communicate back and forth sending and receiving messages on the message queue. I learned how to use message queues, and what the benefits of message queues are over previous IPC methods like Shared Memory and Pipes. I learned how message queues have synchronization provided so there doesn't need to be a flag variable, but the IPC\_NOWAIT flag can be set if a process wants to return immediately, or the flag can be 0 to wait until a message is able to be sent or received. This makes the code easier to write, but the cost is slower access compared to Shared Memory. But, since each message is separate from other messages, it can provide safety for values compared to shared values in Shared Memory.

## Source Code

```
// processA.c
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/msg.h>
```

```
#include <string.h>
#include <signal.h>

#include "msg.h"

void INThandler(int);
int msid;

void closeMessageQueue(int msqid) {
    // close message queue
    if (msgctl(msid, IPC_RMID, NULL) == -1) {
        // error closing message queue
        fprintf(stderr, "Failed to close message queue\n");
    }
    else {
        printf("Closed Message Queue Successfully\n");
        fflush(stdout);
    }
}

int main(void) {
    signal(SIGINT, INThandler);

    // processA creates a new message queue with permissions rw-rw-rw
    msid = msgget(MSG_QUEUE_KEY, IPC_CREAT | 0666);
    if (msid == -1) {
        // error creating new message queue
        fprintf(stderr, "Failed to create new message queue\n");
        exit(1);
    }

    // we don't have to worry about synchronization since message queues
    handle that in the msgsnd and msgrcv functions

    // set processA's to send type 1, and receive type 2
```

```
long sendMsgType = 1;
long receiveMsgType = 2;

struct msg sendMsg;
struct msg receiveMsg;
// initialize defaults (will be using sendMsg and receiveMsg buffers
// so mtype will persist until changed manually)
sendMsg.mtype = sendMsgType;
receiveMsg.mtype = receiveMsgType;

while (1) {
    printf("A > ");
    fflush(stdout);
    // get user input and save to sendMsg
    fgets(sendMsg.mtext, MSG_DATA_LENGTH, stdin);
    // remove \n character than fgets keeps
    sendMsg.mtext[strlen(sendMsg.mtext) - 1] = '\0';

    // send message to Queue
    if (msgsnd(msid, &sendMsg, MSG_DATA_LENGTH, IPC_NOWAIT) == -1) {
        // error
        fprintf(stderr, "Failed to send message: type %ld, text
'%s'\n", sendMsg.mtype, sendMsg.mtext);
        closeMessageQueue(msid);
        exit(1);
    }
    else {
        // sent successful
        // If processA sent "Exit" to processB, then break
        if (strncmp(sendMsg.mtext, "Exit", strlen(sendMsg.mtext)) ==
0) {
            break;
        }
    }

    printf("Waiting ...\n");
```



```
    fflush(stdout);
    // flag is 0, so it will wait until it receives a message
    if (msgrcv(msid, &receiveMsg, MSG_DATA_LENGTH, receiveMsgType,
0) == -1) {
        // error
        fprintf(stderr, "Failed to receive message: type %ld\n",
receiveMsg.mtype);
        closeMessageQueue(msid);
        exit(1);
    }
    else {
        // receive successful

        // if B sent "Exit" to A, exit A
        if (strncmp(receiveMsg.mtext, "Exit",
strlen(receiveMsg.mtext)) == 0) {
            break;
        }

        printf("B: '%s'\n", receiveMsg.mtext);
        fflush(stdout);
    }
}

closeMessageQueue(msid);
exit(0);
}

void INThandler(int sig) {
    closeMessageQueue(msid);
    exit(0);
}
```

```
// processB.c
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/msg.h>

#include <string.h>
#include <signal.h>

#include "msg.h"

void INThandler(int);
int msid;

int main(void) {
    signal(SIGINT, INThandler);

    // processB looks for existing message queue
    msid = msgget(MSG_QUEUE_KEY, 0);
    if (msid == -1) {
        // error getting message queue
        fprintf(stderr, "Failed to get existing message queue\n");
        exit(1);
    }

    // we don't have to worry about synchronization since message queues
    handle that in the msgsnd and msgrcv functions

    // set processB's to send type 2, and receive type 1
    long sendMsgType = 2;
    long receiveMsgType = 1;

    struct msg sendMsg;
    struct msg receiveMsg;
```

```
// initialize defaults (will be using sendMsg and receiveMsg buffers
so mtype will persist until changed manually)
sendMsg.mtype = sendMsgType;
receiveMsg.mtype = receiveMsgType;

while (1) {
    printf("Waiting ...\n");
    fflush(stdout);
    // flag is 0, so it will wait until it receives a message
    if (msgrcv(msid, &receiveMsg, MSG_DATA_LENGTH, receiveMsgType,
0) == -1) {
        // error
        fprintf(stderr, "Failed to receive message: type %ld\n",
receiveMsg.mtype);
        exit(1);
    }
    else {
        // read successful

        // if received message text is "Exit", then exit processB
        if (strncmp(receiveMsg.mtext, "Exit",
strlen(receiveMsg.mtext)) == 0) {
            break;
        }

        printf("A: '%s'\n", receiveMsg.mtext);
        fflush(stdout);
    }

    printf("B > ");
    fflush(stdout);
    // get user input and save to sendMsg
    fgets(sendMsg.mtext, MSG_DATA_LENGTH, stdin);
    // remove \n character than fgets keeps
    (sendMsg.mtext)[strlen(sendMsg.mtext) - 1] = '\0';
```

```

        // send message to Queue
        if (msgsnd(msid, &sendMsg, MSG_DATA_LENGTH, IPC_NOWAIT) == -1) {
            // error
            fprintf(stderr, "Failed to send message: type %ld, text
%s'\n", sendMsg.mtype, sendMsg.mtext);
            exit(1);
        }
        else {
            // send successful

            // if processB sent "Exit" to processA, exit processB
            if (strncmp(sendMsg.mtext, "Exit", strlen(sendMsg.mtext)) ==
0) {
                break;
            }
        }
    }

    exit(0);
}

void INThandler(int sig) {
    exit(0);
}

```

```

//msg.h
#ifndef MSG_H
#define MSG_H

#include <sys/types.h>

#define MSG_DATA_LENGTH 100

```

```
struct msg
{
    __syscall_slong_t mtype;    /* type of received/sent message */
    char mtext[MSG_DATA_LENGTH]; /* text of the message */
};

key_t MSG_QUEUE_KEY = 12345;

#endif
```